

UNIVERSITY OF ZAGREB
FACULTY OF ELECTRICAL ENGINEERING AND COMPUTING

MASTER THESIS No.

**EXTENDING OSTRACK WITH GRU-BASED
MEMORY TOKENS FOR VISUAL OBJECT
TRACKING**

Student's name and surname

Zagreb, June 2026

SVEUČILIŠTE U ZAGREBU
FAKULTET ELEKTROTEHNIKE I RAČUNARSTVA

DIPLOMSKI RAD br.

**EXTENDING OSTRACK WITH GRU-BASED
MEMORY TOKENS FOR VISUAL OBJECT
TRACKING**

Student's name and surname

Zagreb, June 2026

Acknowledgements

Contents

Acknowledgements	i
Introduction	1
1 Background and Related Work	3
1.1 Single-object tracking	3
1.2 From Siamese trackers to transformer trackers	4
1.3 Template and Search Crop Notation	6
1.4 Evaluation metrics and common benchmarks	7
1.5 Motivation for temporal memory	8
2 OTrack Baseline	10
2.1 OTrack Tracker	10
2.2 ViT Backbone	11
2.3 Candidate elimination	15
3 Proposed Memory-Augmented OTrack	19
3.1 Memory-token design motivation	19
3.2 Direct transformer-based memory update	19
3.3 Limitations of direct memory update	21
3.4 Variant B: GRU-based memory update and two-stage architecture	23
4 Sequence-Based and Inference-Like Training	26
4.1 Sequence-based training for recurrent memory	26
4.2 Dynamic search cropping	27
4.3 Truncated Backpropagation Through Time in the training loop	28

5	Auxiliary Memory Supervision	30
5.1	Auxiliary Loss Concept	30
5.2	Sampled Discriminative Filter Loss	30
5.3	DiMP Target-Match Loss	32
5.4	Template Blurring	34
6	Experimental Setup	36
6.1	Dataset and protocol	36
6.2	Model configuration	36
6.3	Hyperparameter and optimizer considerations	37
6.4	Evaluation protocol	37
7	Results and Discussion	39
7.1	Quantitative results	39
7.2	Training Dynamics	40
7.3	Computational Cost and Token Budget	42
7.4	Limitations	43
7.5	Future work	44
	Conclusion	45
	References	46
	Abstract	49
	Summary	50
	Abbreviations	52

Introduction

Visual object tracking is the task of locating a target object throughout a video after the target has been specified in the first frame. The tracker receives an initial bounding box and then predicts the target position in each subsequent frame. Although the input definition is simple, the task is challenging because the target can move, change scale, deform, become partially occluded, or appear in a cluttered background. A tracker must therefore combine the initial target description with current visual evidence and, ideally, with temporal information collected during the sequence.

This thesis investigates a memory extension of OTrack, a one-stream transformer tracker. OTrack is a suitable baseline because it jointly processes template and search tokens in a single transformer backbone. The goal of this work is not to replace the whole tracker, but to test whether explicit recurrent memory tokens can be inserted into the OTrack token stream and trained on video sequences in a meaningful way.

The proposed experimental model is called MemOTrack. It extends the original template-search token sequence with memory tokens and updates these tokens using a two-stage gated recurrent unit (GRU) mechanism. The update combines temporal information from previous frames with layer-wise memory refinement. This design is intended to give the tracker an explicit recurrent state while preserving the main structure of the OTrack baseline.

The work is not limited to an architectural change. Independent template-search frame pairs are insufficient for training the recurrent memory state in a realistic way, because the memory state should evolve over an ordered sequence of frames. For this reason, the sampling pipeline was first rewritten to return video frames as ordered sequences. It was later extended with dynamic search cropping, where the next search

crop is generated from the previous prediction rather than from the current ground-truth target position. This final version better matches the inference setting, where the tracker does not know the target box after initialization.

The thesis also studies a practical problem that appears when adding memory to a strong baseline: the new memory branch can be ignored. Auxiliary memory loss heads were explored to make memory tokens predictive, and template blurring was considered as a regularization strategy that weakens over-reliance on the clean initial template. These ideas do not guarantee improvement, but they broaden the study from an architectural modification to the training and supervision conditions required for a memory pathway to carry useful temporal information.

The main contributions of the thesis are:

- introduction of recurrent memory tokens into the OTrack ViT backbone;
- implementation of a two-stage GRU memory update;
- rewriting of the training sampler for ordered video sequences;
- implementation of inference-like dynamic search cropping;
- exploration of memory loss heads and template blurring;
- experimental comparison with CE-enabled OTrack baselines at 256 and 384 search resolutions.

1 Background and Related Work

1.1 Single-object tracking

Single-object tracking is a sequential computer vision problem. In the first frame of a video, the target is identified by a bounding box. In later frames, the tracker predicts the target location without receiving additional manual annotation. This distinguishes tracking from ordinary object detection: the class of the object is not enough, because the tracker must follow a particular instance. If several similar objects are present, the correct output is the original target, not merely any object of the same category.

Given frames $I_1, I_2, \dots, I_T$ and the initial box B_1 , a tracker predicts boxes $\hat{B}_t$ for $t > 1$.

The tracker may also maintain an internal state. A memory-free tracker can be written as a function of the template and the current search region, while a recurrent tracker additionally uses a state carried from previous frames. This distinction is central to this thesis, because MemOSTrack explicitly introduces recurrent memory tokens into the OS-Track backbone.

Tracking is difficult because the visual appearance of the target is not fixed. The target may rotate, deform, change illumination, become blurred by motion, or be partly hidden. The background can contain objects that look similar to the target. The camera may move, producing large apparent target displacement between frames. A tracking method must therefore balance stability and adaptability: it should not drift to a distractor, but it must adapt when the true target changes appearance.

Single-object tracking (SOT) is the task of localizing the same target object throughout a video sequence, given only its bounding box in the first frame. Early tracking methods were mainly based on hand-crafted appearance and motion models, including

mean-shift tracking [1], optical flow-based methods [2], particle filters [3], and template matching approaches [4]. These methods were often efficient and conceptually simple, but they struggled under large appearance changes, occlusion, background clutter, and scale variation. Around the early 2010s, discriminative correlation filter trackers became one of the dominant families in visual tracking. Methods such as MOSSE [5], CSK [6], KCF [7], and DSST [8] learned online filters that could localize the target efficiently in the frequency domain, while DSST also introduced a separate scale estimation mechanism. Later, more advanced correlation-filter trackers such as ECO [9] improved robustness by using more compact model updates and stronger feature representations. This period is commonly treated as the correlation-filter era of SOT, before deep Siamese and transformer-based trackers became the dominant direction.

Table 1.1 - Common tracking challenges.

Challenge	Example	Why it is difficult
Occlusion	Target is partially hidden	The visible evidence becomes incomplete and the tracker may drift.
Scale change	Target approaches the camera	The size of the target box changes over time.
Deformation	Person, animal, or flexible object changes shape	Appearance no longer matches the initial template exactly.
Background clutter	Similar objects in the scene	The tracker can confuse target and distractor.
Motion blur	Fast motion or camera movement	Fine details disappear from the search crop.
Out-of-view motion	Target leaves and returns	The tracker state may be corrupted during absence.

1.2 From Siamese trackers to transformer trackers

The state of the art in single-object tracking has largely focused on template-to-target matching, where the tracker relies on visual cues from two image regions: a template crop that describes the target appearance and a search crop in which the target has to be

localized. A major step in this direction was introduced by SiamFC [10], which formulated tracking as a fully convolutional Siamese matching problem between an exemplar image and a larger search image. Later, SiamRPN [11] extended this idea by combining Siamese feature extraction with a region proposal subnetwork, allowing the tracker to perform target classification and bounding-box regression. In these early Siamese trackers, the template and search images are processed by two branches with shared convolutional weights, after which the extracted features are compared or fused to estimate the target location.

Later trackers introduced transformer-based components for richer relation modeling between the target and the search region. STARK [12] used an encoder-decoder transformer to model global spatio-temporal dependencies and to directly predict the target bounding box. OSTRack [13] further simplified the tracking pipeline by introducing a one-stream transformer framework. Instead of extracting template and search features independently and then applying a separate relation module, OSTRack jointly processes template and search tokens inside the same backbone, so that feature extraction and relation modeling are unified. Recent high-performing trackers on GOT-10k include temporal and autoregressive models such as ODTrack [15], ARTrackV2 [16], and ARPTrack [17]. These models indicate that recent tracking research increasingly explores stronger temporal modeling, token propagation, and sequence-level pretraining.

All of these models share a similar high-level tracking structure. They accept a template crop, which represents how the target looks, and a search crop, which contains the region of the current frame where the target is expected to appear. The tracker then extracts visual features from these two inputs and estimates the target position by comparing, fusing, or jointly modeling the template and search representations. The main architectural differences lie in how these features are extracted and how the relation between the template and search region is modeled: earlier Siamese trackers rely mainly on convolution and cross-correlation, while transformer-based trackers use attention mechanisms to allow richer interaction between template and search features.

It is also important to note that feature extraction in modern trackers is often based on pretrained backbones. These backbones are trained on large-scale image or video datasets before being adapted to tracking. This is important because tracking datasets

alone are usually not sufficient for learning strong low-level and semantic visual features from scratch. In practice, the pretrained backbone provides general visual representations, while tracker training teaches the model how to use these representations for target localization and target-background discrimination.

Another important family of trackers is represented by DiMP [18]. Unlike purely Siamese approaches that mainly compare a fixed template with the current search region, DiMP learns a discriminative target model during tracking. Its main idea is to predict a target-specific classifier that can distinguish the target from the background using information available from the video sequence. This makes DiMP especially relevant to this thesis because it explicitly uses sequence information to adapt the target model online, instead of relying only on a static initial template. The later PrDiMP [19] extends this direction with a probabilistic regression formulation for visual tracking. Although DiMP-style trackers are architecturally different from OTrack, they show that online adaptation and sequence-level information can improve tracking robustness. This motivates the idea of adding recurrent memory to a template-search transformer tracker.

1.3 Template and Search Crop Notation

This thesis follows the standard template-search notation used in many modern single-object trackers. The **template crop** is denoted by T . It is extracted from the first annotated frame and contains the target object we are tracking. It is used to provide the visual description of the object that must be tracked.

The **search crop** at time t is denoted by S_t . It is extracted from the current frame around the expected target location. Usually, a neural-network tracker does not process the full image directly. Instead, it receives a smaller resized crop around the expected target location. This location is usually estimated from the bounding-box prediction in the previous frame.

There are two main reasons why a search crop is used instead of the whole image. First, it reduces computation by allowing the model to process only the region where the target is expected to appear. This is especially important for transformer-based trackers, because transformer-style attention is quadratic in the number of tokens [21]. For exam-

ple, with 16×16 patches, a 256×256 search crop contains 256 tokens, whereas a full HD frame would contain more than 8000 tokens. The corresponding attention matrix is therefore much smaller for the search crop than for the full image.

Second, the search crop provides a useful spatial prior. Since the crop is usually centered around the previous target bounding box, it gives the tracker information about where the target was located in the previous frame and what its approximate scale was. In short-term tracking, the target is usually expected to remain near its previous location, so a local search region is often sufficient. The tracker then predicts the new target box inside this resized crop, and the prediction is converted back to the coordinate system of the original image.

In many trackers, the template is cropped only at initialization and remains fixed during the sequence. Other trackers update or maintain an additional target representation during inference. FEAR [14], for example, uses a dual-template representation for object model adaptation, where a dynamic template complements the initial template during tracking.

In contrast, trackers such as OTrack store only the initial template extracted from the first frame and do not maintain an explicit updated visual template during inference. This makes the tracker dependent on the initial target appearance. If the target rotates, deforms, changes illumination, or becomes partially occluded, the initial template may no longer describe the current appearance well. The proposed memory mechanism is designed to reduce this limitation. MemOTrack adds latent recurrent memory tokens inside the transformer token sequence. These memory tokens are intended to store additional temporal cues about the changing target appearance.

1.4 Evaluation metrics and common benchmarks

Because single-object tracking is a widely studied problem, it has established benchmarks and evaluation metrics. These benchmarks make it possible to compare trackers under a common protocol, using the same test sequences and the same localization measures.

In this thesis, the main benchmark is GOT-10k. GOT-10k is a generic object tracking

benchmark designed to evaluate how well a tracker generalizes to unseen object classes [20]. It follows a one-shot tracking protocol: the target is specified by a bounding box in the first frame, and the tracker must localize the same object in the remaining frames. The test annotations are hidden, so final results are obtained through the official evaluation server.

The standard localization measure used in these evaluations is intersection over union (IoU), also called overlap. For a predicted box and a ground-truth box, IoU is the area of their intersection divided by the area of their union. A value of 1 means perfect overlap, while a value of 0 means that the boxes do not overlap.

$$\text{IoU}(B_{\text{pred}}, B_{\text{gt}}) = \frac{|B_{\text{pred}} \cap B_{\text{gt}}|}{|B_{\text{pred}} \cup B_{\text{gt}}|}$$

GOT-10k reports Average Overlap (AO), SR0.50, and SR0.75 [20]. AO is computed as the average IoU over all evaluated frames:

$$\text{AO} = \frac{1}{N} \sum_{i=1}^N \text{IoU}(B_{\text{pred}}^i, B_{\text{gt}}^i).$$

SR0.50 is the proportion of frames whose overlap is at least 0.50, while SR0.75 is the proportion of frames whose overlap is at least 0.75. SR0.75 is stricter and is more sensitive to precise localization errors. These metrics are useful together because a tracker may localize the target approximately while still producing less accurate bounding boxes.

1.5 Motivation for temporal memory

The motivation for adding temporal memory to a tracker comes from several limitations of using only the initial template and the current search crop. First, the initial template becomes less reliable as the video progresses. It captures the target only in the first frame, while the target may later rotate, deform, change illumination, change scale, or become temporarily occluded. The current search crop can also be unreliable: the object may be blurred, partly hidden, or surrounded by visually similar distractors. In such cases, a tracker needs a mechanism that can dynamically encode the most relevant appearance

cues from previous frames instead of relying only on the first template.

Second, tracking is inherently temporal. A simple thought experiment makes this clear. Imagine looking at a single frame of grass and trying to identify a camouflaged snake. In that one image, the snake may be almost indistinguishable from the background. However, once the video starts playing, coherent movement across consecutive frames can make the same object visible. The object has not necessarily changed its appearance, but the temporal context makes it easier to separate from the background. Additionally, a temporal memory can potentially encode motion cues. A memory state does not need to store only appearance; it can also learn information about recent target movement, direction, speed, and stable object identity across frames. This can be useful when the current frame is ambiguous, because the previous trajectory and recent observations can provide additional evidence about where the target is likely to be.

A simple way to use temporal information would be to provide the transformer with tokens from previous frames together with the current search crop. However, this quickly leads to a temporal token explosion. For example, if one search crop contains 256 tokens, then 20 previous search crops already contain 5120 tokens before adding template tokens. Even if previous-frame key-value representations are cached, the current frame must still attend to a growing set of stored tokens, increasing both memory usage and attention cost.

Therefore, instead of storing all previous visual tokens, an efficient tracker requires a reliable mechanism for dynamically maintaining a compressed hidden state. Such a state should preserve useful temporal cues from previous frames while keeping computation bounded as the sequence length increases.

2 OTrack Baseline

2.1 OTrack Tracker

OTrack is a real-time single-object visual tracker introduced by Ye et al. in *Joint Feature Learning and Relation Modeling for Tracking: A One-Stream Framework*, accepted at ECCV 2022 [13]. It is used in this thesis as the baseline tracker because it provides a strong transformer-based architecture while keeping the tracking pipeline relatively simple.

During inference, OTrack processes one search frame at a time. It receives a fixed template crop from the first annotated frame and a search crop from the current frame. The tracker localizes the target inside the search crop, and the predicted bounding box is then used to define the search region for the next frame.

The main architectural idea of OTrack is its one-stream design. Instead of extracting template and search features separately and then applying a separate relation module, OTrack converts both crops into tokens and processes them jointly in a single Vision Transformer (ViT) backbone. As a result, feature extraction and template-search relation modeling are performed inside the same backbone. Self-attention over the combined token sequence allows template and search tokens to interact throughout the network, which helps the model learn target-oriented search features [13].

After the ViT backbone, OTrack uses a lightweight fully convolutional prediction head. The final search tokens are reshaped into a two-dimensional spatial feature map and passed to this head. The head has three fully convolutional output branches, which predict a classification score map, a local offset map, and a normalized bounding-box size map. The position with the highest classification score is selected as the target position. The backbone is initialized from a pretrained ViT model.

OSTrack also supports candidate elimination, which progressively removes less relevant search tokens to reduce computation; this mechanism is described separately in Section 2.3. For this thesis, the most important property is that OSTRack has strong template-search interaction, but no explicit recurrent memory state across frames.

2.2 ViT Backbone

OSTrack is built on a pretrained Vision Transformer (ViT) backbone. Most of the modifications in this thesis are focused around the ViT backbone. ViT is a powerful image-processing architecture that has shown strong performance in computer vision and, with large-scale pretraining, can match or outperform state-of-the-art convolutional networks on several image-recognition benchmarks. ViT represents an image as a sequence of patch tokens and processes this sequence using transformer encoder layers [22]. In OSTRack, this structure is adapted to tracking by applying patch embedding to both the template crop and the search crop, and then processing the resulting tokens jointly in the same backbone [13].

During the ViT processing of an image crop X , the image is first divided into non-overlapping patches:

$$N = \frac{HW}{P^2}, \quad X_p = [x_1, x_2, \dots, x_N]^\top \in \mathbb{R}^{N \times P^2 C}. \quad (2.1)$$

where:

- $X \in \mathbb{R}^{H \times W \times C}$ is the input image crop;
- H and W are the height and width of the crop;
- C is the number of image channels;
- P is the patch size, assuming that H and W are divisible by P ;
- N is the number of patches;
- $x_i \in \mathbb{R}^{P^2 C}$ is the flattened vector of the i -th patch;

- X_p is the matrix containing all flattened patch vectors.

The flattened patches are projected into a D -dimensional embedding space, and positional embeddings are added:

$$E_X = X_p W_E + b_E, \quad Z_0 = E_X + P_{\text{pos}}. \quad (2.2)$$

where:

- $E_X \in \mathbb{R}^{N \times D}$ is the patch-embedding sequence;
- $W_E \in \mathbb{R}^{P^2 C \times D}$ is the learnable patch-projection matrix;
- $b_E \in \mathbb{R}^D$ is the learnable projection bias, broadcast across all patches;
- D is the token embedding dimension;
- $P_{\text{pos}} \in \mathbb{R}^{N \times D}$ is the positional embedding;
- $Z_0 \in \mathbb{R}^{N \times D}$ is the initial token sequence passed to the transformer encoder.

A standard pre-normalization ViT encoder layer first applies Layer Normalization and Multi-Head Self-Attention, followed by a residual connection:

$$A_l = \text{LN}(Z_{l-1}), \quad U_l = Z_{l-1} + \text{MSA}(A_l). \quad (2.3)$$

where:

- l is the transformer layer index;
- Z_{l-1} is the input token sequence to layer l ;
- $\text{LN}(\cdot)$ denotes Layer Normalization;
- A_l is the normalized token sequence;
- $\text{MSA}(\cdot)$ denotes Multi-Head Self-Attention;

- U_l is the token sequence after self-attention and residual addition.

In the CE-enabled OTrack backbone, candidate elimination is inserted after this attention-and-residual step:

$$\bar{U}_l = \begin{cases} \text{CE}_l(U_l), & \text{if CE pruning is enabled at layer } l, \\ U_l, & \text{otherwise.} \end{cases} \quad (2.4)$$

where:

- $\text{CE}_l(\cdot)$ denotes the candidate elimination operation at layer l ;
- $\bar{U}_l$ is the token sequence after optional CE pruning;
- if CE is not enabled at layer l , the sequence is left unchanged.

Candidate elimination is specific to OTrack and is only present in CE-enabled configurations. It is not part of the original ViT encoder block. The detailed CE scoring and pruning procedure is described separately in Section 2.3.

After the optional CE step, the feed-forward MLP sublayer is applied with another residual connection:

$$Z_l = \bar{U}_l + \text{MLP}(\text{LN}(\bar{U}_l)). \quad (2.5)$$

where:

- $\bar{U}_l$ is the input to the MLP part of the encoder layer;
- $\text{MLP}(\cdot)$ is the feed-forward network;
- Z_l is the output token sequence of layer l .

The Multi-Head Self-Attention operation lets each token aggregate information from other tokens in the same sequence. For one attention head h , it is written as:

$$\begin{aligned}
Q_h &= AW_Q^h, & K_h &= AW_K^h, & V_h &= AW_V^h, \\
\text{head}_h(A) &= \text{softmax} \left(\frac{Q_h K_h^\top}{\sqrt{d_h}} \right) V_h, \\
\text{MSA}(A) &= \text{Concat}(\text{head}_1(A), \dots, \text{head}_{n_h}(A)) W_O.
\end{aligned} \tag{2.6}$$

where:

- A is the normalized input sequence to the attention operation;
- Q_h, K_h , and V_h are the query, key, and value matrices for head h ;
- W_Q^h, W_K^h , and W_V^h are learnable projection matrices for head h ;
- d_h is the feature dimension of one attention head;
- n_h is the number of attention heads;
- W_O is the output projection matrix after concatenating all heads.

The MLP is applied independently to each token:

$$\text{MLP}(z) = W_2 \text{GELU}(W_1 z + b_1) + b_2. \tag{2.7}$$

where:

- $z \in \mathbb{R}^D$ is one token vector;
- W_1, W_2, b_1 , and b_2 are learnable MLP parameters;
- $\text{GELU}(\cdot)$ is the activation function used inside the feed-forward block.

OTrack applies this ViT backbone to a pair of tracking crops. Let T denote the template crop and S_t denote the search crop at frame t . Their token sequences are

$$\begin{aligned}
T_0 &= \text{PatchEmbed}(T) + P_T, & S_{t,0} &= \text{PatchEmbed}(S_t) + P_S, \\
Z_{t,0}^{\text{track}} &= [T_0, S_{t,0}].
\end{aligned} \tag{2.8}$$

where:

- T is the template crop from the first annotated frame;
- S_t is the search crop from frame t ;
- $\text{PatchEmbed}(\cdot)$ denotes patch flattening followed by linear projection;
- P_T and P_S are positional embeddings for the template and search tokens;
- $T_0 \in \mathbb{R}^{N_T \times D}$ is the initial template-token sequence;
- $S_{t,0} \in \mathbb{R}^{N_S \times D}$ is the initial search-token sequence;
- $[\cdot, \cdot]$ denotes concatenation along the token dimension;
- $Z_{t,0}^{\text{track}}$ is the combined OTrack input sequence for frame t .

After L transformer layers, the backbone output is

$$Z_{t,L}^{\text{track}} = \text{ViTBackbone}(Z_{t,0}^{\text{track}}) = [T_L^{(t)}, S_{t,L}]. \quad (2.9)$$


where:

- L is the number of transformer layers;
- $Z_{t,L}^{\text{track}}$ is the final backbone output for frame t ;
- $T_L^{(t)}$ denotes the processed template-token representation while tracking frame t ;
- $S_{t,L}$ denotes the processed search-token representation for frame t .

The final search-token part $S_{t,L}$ is reshaped into a spatial feature map and passed to the prediction head.

2.3 Candidate elimination

OTrack introduces candidate elimination (CE) to reduce the number of search tokens processed by the transformer backbone. In a template-search tracker, the search crop



contains both the target and many background patches. CE removes search tokens that are unlikely to correspond to the target, reducing computation while keeping the most relevant candidates [13].

OTrack chooses candidates using attention between the template and search tokens. After the multi-head attention operation in a transformer block, a representative template token is used to score the search tokens. In the standard setting, this representative token is the center template token, because the target is normally centered in the template crop. If this token is denoted by ϕ and the transformer has M attention heads, the score of search token x is computed by averaging attention from ϕ to x over all heads:

$$w_x^\phi = \frac{1}{M} \sum_{m=1}^M w_x^\phi(m).$$

Search tokens with the highest scores are kept, while lower-scoring search tokens are removed from the sequence for later transformer layers. If k tokens are kept from n current search tokens, the keep ratio is:

$$\rho = \frac{k}{n}.$$

The kept search-token indices can be written as:

$$I_{\text{keep}} = \text{TopK}(w_x^\phi, k).$$

The original order and positions of the remaining search tokens are stored. Before the prediction head, the kept search tokens are restored to their spatial order and the removed positions are zero-padded so that the spatial search feature map can be reconstructed [13].

This matters for MemOTrack because the proposed model adds memory tokens to the OTrack token sequence. In this implementation, CE is applied only to search tokens. Template tokens are preserved, and memory tokens are also preserved because they are recurrent latent states rather than spatial search candidates. Memory tokens participate

in attention, but they are not used to score search candidates and are not removed by CE. Thus, candidate elimination keeps its original role of pruning search-region candidates, while the recurrent memory stream remains available throughout the backbone.

3 Proposed Memory-Augmented OTrack

3.1 Memory-token design motivation

MemOTrack extends OTrack by adding explicit memory tokens to the transformer backbone. These memory tokens are inserted into the same token sequence as the template and search image feature embeddings, so they participate in ViT attention together with the visual tokens. However, unlike template and search tokens, memory tokens are latent hidden-state vectors rather than image patches, and they are carried across frames during tracking.

This design allows the ViT backbone to continuously read from and refine the memory representation through attention. Visual tokens can interact with memory tokens while the current template and search features are processed, and the resulting memory representation can store information that may be useful in later frames.

The architectural changes are mainly applied inside the backbone. The patch embedding and OTrack prediction head are kept close to the baseline implementation, so the model remains an OTrack-style tracker with an added recurrent memory stream.

3.2 Direct transformer-based memory update

The first architectural variant adds memory tokens to the OTrack backbone and lets the ViT layers update them directly. In this version, memory tokens are concatenated with the template and search tokens and are processed by the same transformer layers as the visual tokens. The goal is to test whether the existing transformer operations, including self-attention, residual connections, and feed-forward layers, are sufficient to learn a useful memory update strategy.

In the direct transformer-based update, each ViT layer produces transformed search, template, and memory tokens in a single forward pass:

$$[S_{t,l}, T_{t,l}, M_{t,l}] = \text{ViTLayer}_l([S_{t,l-1}, T_{t,l-1}, M_{t,l-1}]).$$

Here:

- $S_{t,l}$ denotes the search tokens after layer l while processing frame t ;
- $T_{t,l}$ denotes the template tokens after layer l while processing frame t ;
- $M_{t,l-1}$ is the memory input from the previous transformer layer in the current frame;
- $M_{t,l}$ is the memory output produced directly by the ViT layer.

All token groups have the same embedding dimension D , so they can be concatenated along the token dimension and processed by the same transformer block.

For the first frame of a sequence, no previous memory state exists. The model therefore uses a learned initial memory state. Then, after each frame has been processed, the backbone returns an updated memory state in the auxiliary output. This updated memory state is then provided to the backbone when processing the next search frame, together with the template and the next search crop.

To distinguish memory tokens from visual tokens, the model adds a learned memory-token type embedding to the memory tokens before they are processed by the backbone. This embedding shifts the memory tokens in feature space and gives the transformer a learned signal that these prefix tokens represent memory states rather than image patches.

In the CE-enabled backbone, memory tokens are intentionally excluded from candidate-elimination decisions. First, they are not pruned, because CE removes only search tokens. Second, they are excluded from the template mask used for CE scoring, so they do not directly influence which search tokens are kept or removed. Thus, CE pruning is decided without treating them as either removable candidates or scoring tokens, so memory tokens can influence it only indirectly through their participation in

transformer attention.




Figure 3.1 - Direct transformer-based memory update. The template and search images are first converted into visual tokens by tokenization and linear projection. A learned initial memory state, or the memory state returned from the previous frame, is inserted into the transformer token sequence. Each ViT layer jointly processes search, template, and memory tokens through self-attention. The updated memory state is passed through the backbone and returned for use in the next frame, while the final visual tokens are reshaped and passed to the OTrack prediction head.

3.3 Limitations of direct memory update

The direct-update variant exposes two fundamental limitations regarding long-term temporal modeling. While the standard Vision Transformer architecture is effective for spatial feature extraction and template-search matching, relying on it alone for recurrent memory updates is insufficient.

The first limitation is the lack of explicit gating for memory preservation. In the direct variant, memory tokens are updated by the same transformer block as the visual tokens. For a standard pre-normalization transformer layer, the output of the attention sublayer

is defined by a residual connection:

$$X'_l = X_{l-1} + \text{MSA}(\text{LN}(X_{l-1})).$$

Here, X_{l-1} represents the input token sequence, including search, template, and memory tokens. $\text{LN}(\cdot)$ denotes Layer Normalization, and $\text{MSA}(\cdot)$ denotes Multi-Head Self-Attention.

While the attention mechanism is effective for determining which visual tokens should influence the memory state, it does not provide a dedicated mechanism for controlling how much of the existing memory should be preserved. The residual connection described by the equation above can carry the previous token representation forward. However, it adds the previous token representation and the new attention update without explicitly deciding how much each should contribute. Over long tracking sequences of tens or hundreds of frames, the lack of explicit gating can become a limitation. In the direct-update variant, repeated ungated updates may gradually make the memory state drift or become less representative of the target.

The second limitation is the susceptibility to gradient instability during Backpropagation Through Time due to an absence of temporal skip connections. In recurrent neural networks, there are often temporal skip connections between the memory representations of the same layer across time. However, in the direct-update design, recurrence is present only in the form of a feedback loop, where the last layer of the current frame is connected to the first layer of the next frame. This means that when unrolling the memory tokens over T frames through an L -layer backbone, the gradient has to travel through all layers of the backbone for each frame:

$$\frac{\partial L_T}{\partial M_1} \supset \frac{\partial L_T}{\partial M_T} \prod_{t=2}^T \prod_{l=1}^L \frac{\partial M_{t,l}}{\partial M_{t,l-1}}.$$

Here, L_T is the loss at frame T , M_1 and M_T are the memory states at the first and last frames, and $\frac{\partial M_{t,l}}{\partial M_{t,l-1}}$ represents the local Jacobian matrix at layer l of frame t . This expression isolates the recurrent memory path; the full training gradient also contains

paths through search tokens, template tokens, the prediction head, and the backbone parameters.

Because the search and template tokens do not propagate their hidden states across frames, the recurrent temporal path is carried by the memory tokens, represented here by $\frac{\partial M_{t,l}}{\partial M_{t,l-1}}$. Suppose we have a 12-layer transformer backbone over a 20-frame sequence; this chain expands to 240 sequential Jacobian matrix multiplications. This deep computational graph can contribute to error accumulation, causing gradient vanishing or explosion. One way to address this issue is to add explicit temporal skip connections.

To address both memory drift and gradient instability, a more intentional architectural design is required. Specifically, the architecture must decouple the *proposal* of new memory from the actual *update* of the memory state, and it must introduce explicit temporal skip connections to stabilize gradient flow. This motivates the second architectural variant, which replaces the direct update with a two-way gated recurrent mechanism, as detailed in the following section.

3.4 Variant B: GRU-based memory update and two-stage architecture

To implement this controlled update, the second architectural variant introduces a Gated Recurrent Unit (GRU) mechanism to the memory stream. A GRU is a recurrent module that updates a hidden state using learned gates [23]. These gates control how much of the previous state is preserved and how much new information is accepted. This is highly suitable for memory tokens because the tracker must retain reliable target information while still adapting when the target appearance changes.

In this version, the transformer layer does not directly produce the final memory. Instead, the transformer output is treated as a candidate memory representation:

$$[S_{t,l}, T_{t,l}, M'_{t,l}] = \text{ViTLayer}_l([S_{t,l-1}, T_{t,l-1}, M_{t,l-1}]).$$

The value $M'_{t,l}$ is the memory suggested by attention at the current frame and layer. It is not accepted directly. Instead, it is passed to a two-stage GRU update. The first GRU

combines the candidate memory with the memory from the previous frame at the same layer, creating an explicit temporal skip connection:

$$\tilde{M}_{t,l} = \text{GRU}_1(M'_{t,l}, M_{t-1,l}).$$

The second GRU combines the intermediate memory with memory from the previous layer in the current frame:

$$M_{t,l} = \text{GRU}_2(\tilde{M}_{t,l}, M_{t,l-1}).$$

The final output of the layer is:

$$Y_{t,l} = [S_{t,l}, T_{t,l}, M_{t,l}].$$

This recurrent update is used only when previous-frame memory is available. On the first frame, the tracker has no temporal memory to compare against, so the GRU update is skipped. The first GRU can be interpreted as temporal fusion. It decides how much memory from frame $t - 1$ should remain when the current candidate memory is introduced. The second GRU can be interpreted as layer-wise fusion. It allows memory from the previous transformer layer to influence the final memory at the current layer, creating a gated depth path for memory tokens.

This design was added late in development. The GRU update makes the memory stream explicitly recurrent and gives the model a learned mechanism for balancing preservation and adaptation.

A related alternative considered during development was a three-way gated update. In that design, the final memory would be a weighted combination of the three available memory sources:

$$M_{t,l} = \alpha_1 M'_{t,l} + \alpha_2 M_{t-1,l} + \alpha_3 M_{t,l-1}, \quad \alpha_1 + \alpha_2 + \alpha_3 = 1.$$




Figure 3.2 - GRU-based memory update with temporal and depth connections inside the backbone.

Such a softmax-style gate would make the three-source structure explicit. The implemented version instead uses two sequential GRU updates. The GRU version is more recurrent in form and uses learned gates inside each update. The three-way gated update remains a useful future architecture to explore because it would treat all three memory sources symmetrically.

4 Sequence-Based and Inference-Like Training

4.1 Sequence-based training for recurrent memory

Standard OTrack training is based on independent template-search pairs. A template crop is sampled from one frame, a search crop is sampled from another frame in the same video, and the model is trained to predict the target box inside the search crop:

$$\text{sample} = (T, S_t, B_t), \quad \hat{B}_t = f(T, S_t).$$

Here, T is the template crop, S_t is the search crop at time t , and B_t is the target box expressed in the coordinate system of the search crop. This setup is sufficient for a memory-free tracker because each training sample can be processed independently. No recurrent state has to be carried from previous search frames.

For MemOTrack, independent pair-based training is insufficient. If the model has a memory state M_t , then this state should be produced from earlier frames in the same video. Resetting memory for every pair prevents the model from learning how memory evolves across time, while carrying memory across unrelated samples would be incorrect.

The sampler was therefore changed to return ordered search-frame sequences [25]. In the causal consecutive setting, the sampled frames follow the same temporal direction as inference:

$$F_t, F_{t+1}, \dots, F_{t+N-1}.$$

The training sample becomes:

$$\text{sample} = (T, S_1, S_2, \dots, S_N, B_1, B_2, \dots, B_N).$$

For each search frame, the model predicts a box and updates its memory state. The memory output from one frame becomes the memory input for the next frame. The important point is that the order is preserved: frame t is processed after frame $t - 1$, not as an independent shuffled sample. This gives the recurrent memory module a meaningful temporal path during training.

4.2 Dynamic search cropping

Sequence-based sampling makes recurrent memory training possible, but it does not by itself remove the mismatch between training and inference. If every search crop is centered using the ground-truth box of the current frame, the tracker receives cleaner inputs than it will receive at test time. The target is usually near the center of the crop, and the model does not experience the consequences of its own previous prediction errors.

During inference, the tracker does not know the current ground-truth target location. The next search crop must be generated from the previous prediction:

$$S_t = \text{crop}(I_t, \hat{B}_{t-1}).$$

Ground-truth-centered training instead uses:

$$S_t = \text{crop}(I_t, B_t^{\text{gt}}).$$

These two settings are fundamentally different. The ground-truth-centered crop uses information that is unavailable during inference, while the prediction-centered crop allows localization errors to influence later inputs. This difference is especially important for recurrent memory because an incorrect crop can also affect the memory state passed to later frames.

The final training pipeline therefore introduced dynamic search cropping. The tracker state is initialized from the first-frame ground-truth box, but later crops are generated from the model’s own previous predictions. At each step, the model receives the template, the current search crop, and the current memory state; it predicts a target box, updates memory, maps the prediction back to image coordinates, and uses that prediction to crop the next frame:

$$S_t = \text{crop}(I_t, \hat{B}_{t-1}), \quad \hat{B}_t, M_t = f(T, S_t, M_{t-1}).$$

This dynamic loop makes training intentionally harder. With ground-truth-centered crops, the model sees clean examples even if its previous predictions would have been inaccurate. With dynamic cropping, an early localization error can move the next crop away from the target center, causing the target to appear near the edge, partially outside the crop, or surrounded by distractors. However, this difficulty is necessary to close the gap between training and inference. If the goal is to train a recurrent memory module, the memory must be exposed to realistic inputs and learn to recover from drift.

Because dynamic cropping makes optimization more difficult, the final training pipeline used a curriculum strategy. During the first 40 epochs, search crops were centered using ground-truth target boxes. After this warm-up phase, inference-like dynamic cropping was enabled, so each next crop was generated from the model’s previous prediction. This allowed the model to first learn stable localization before being exposed to prediction drift.

4.3 Truncated Backpropagation Through Time in the training loop

The introduction of memory tokens transforms the tracker into a recurrent model, necessitating careful management of error backpropagation across temporal iterations. Allowing gradients to flow through an entire video sequence is both computationally expensive and highly demanding on GPU memory. Additionally, unrolling the computational graph over long sequences amplifies the risk of gradient instability, such as vanishing or exploding gradients. A standard approach to address this is Truncated Backpropa-

gation Through Time [27]. In the provided implementation, the model periodically detaches the memory state from the computation graph after a fixed number of frames (specifically, 8 frames during training). This truncation limits how far the gradients can propagate backward in time, balancing the need to learn temporal dependencies with computational feasibility and training stability.

5 Auxiliary Memory Supervision

5.1 Auxiliary Loss Concept

Adding memory tokens to a strong model does not guarantee that the model will use them. OTrack already has a powerful template-search pathway. Since the tracker can solve the training task using template and search tokens alone, the new memory tokens may receive weak or unimportant gradients. This creates the risk that the memory tokens collapse to an uninformative representation or are ignored by the prediction pathway.

This is a common issue when adding auxiliary modules to strong neural networks. The optimization process often follows the dominant low-loss pathway. If the template is clean and the search crop is well-centered, that pathway may rely entirely on the original OTrack features. The memory branch may then have little to no effect on the final prediction.

The thesis explored auxiliary memory supervision. The goal was to force the memory tokens to encode representations that are directly useful for prediction, rather than being present in the token sequence.

5.2 Sampled Discriminative Filter Loss

The first auxiliary loss was designed to explicitly force the memory tokens to encode spatial representations capable of localizing the target. To achieve this, an auxiliary **Memory Head** was introduced to the architecture. This lightweight multi-layer perceptron takes only the recurrent memory tokens as input and projects them into a spatial convolutional filter. Rather than allowing the memory tokens to remain abstract, this approach evaluates their discriminative power by testing how well this predicted filter can localize

the target across the video sequence.

Let f_t denote the memory filter predicted by the memory head at frame t . To evaluate this filter, a random subset of frames, denoted by the index set $\mathcal{S}_t$, is sampled from the current training sequence. For each sampled frame $j \in \mathcal{S}_t$, the filter f_t is applied to the corresponding search region feature map x_j via cross-correlation to produce a spatial response map:

$$s_{t,j} = x_j * f_t.$$

The objective is to minimize the discrepancy between this predicted response map and the ground-truth Gaussian target heatmap y_j . The loss for the filter at frame t is formulated as the Mean Squared Error (MSE) over the sampled subset, combined with an L_2 regularization penalty on the filter weights:

$$L_t = \frac{1}{|\mathcal{S}_t|} \sum_{j \in \mathcal{S}_t} \|x_j * f_t - y_j\|_2^2 + \lambda \|f_t\|_2^2.$$

Here, λ is a hyperparameter controlling the strength of the weight regularization, which prevents the filter values from growing unbounded. The total auxiliary loss is the average of L_t across all frames in the sequence that produce a memory filter.

A critical design choice in this algorithm is the treatment of the search features x_j . During the loss computation, these features are explicitly detached from the computational graph. If the search features were not detached, gradients from this auxiliary loss would flow backward through them and into the ViT backbone for the sampled frames. That would force the backbone to modify its general visual representations to accommodate the memory loss. By detaching the features, the optimization gradients are forced to propagate exclusively backward through the predicted filter f_t , the auxiliary memory head, and into the recurrent memory tokens. This ensures that the loss strictly supervises the memory pathway without inadvertently altering the primary feature extractor.

The primary strength of this loss is that it evaluates the memory representation in a direct, tracking-oriented manner: the memory state must generate a filter that yields a

high response at the target location and a low response on the background, even when applied to frames from different temporal steps.

However, this approach presents two significant limitations. First, because the supervision relies on a randomly sampled subset of frames ($\mathcal{S}_t$), the training signal can exhibit high variance and noise. Second, the computational complexity scales poorly. Evaluating every predicted filter against every other frame in a sequence of length T would require $O(T^2)$ cross-correlation operations. While random sampling reduces this to $O(T \cdot |\mathcal{S}_t|)$, it remains computationally heavy. These limitations motivated the development of a more efficient, sequence-level approach, detailed in the following section.

5.3 DiMP Target-Match Loss

The second auxiliary loss is the DiMP-style target-match loss. Instead of directly testing every predicted memory filter on sampled features from previous frames, this loss first constructs a stronger sequence-level target filter and then trains each predicted memory filter to match it.

This idea is inspired by DiMP [18], where tracking is formulated as learning a target-specific discriminative model from training samples. DiMP showed that online target model prediction can use both target and background information, instead of only matching a target template [18].

The proposed DiMP-style target-match loss is not a reimplementation of the full DiMP tracker. Instead, it borrows the idea of fitting a discriminative target filter from a sequence of search features. The fitted filter is used as a detached pseudo-label for the memory branch. This differs from original DiMP, where the optimized filter is the actual target classifier used for tracking. The simplification is intentional: the goal is not to replace OTrack’s prediction head, but to provide direct supervision that encourages the added memory tokens to encode information useful for target discrimination.

To construct the pseudo-label, the method first collects all available detached search features, $x_1, \dots, x_T$, and their corresponding ground-truth heatmaps, $y_1, \dots, y_T$. An initial filter is formed by taking the mean of the detached memory filters predicted by the auxiliary memory head across the sequence. A sequence-level target filter, f^* , is then fitted

by minimizing the following least-squares objective:

$$L(f) = \text{mean}(\|x * f - y\|_2^2) + \lambda \text{mean}(\|f\|_2^2).$$

Here, f^* is the filter that best matches the whole training sequence under this least-squares objective. The code solves this using a DiMP-style steepest-descent solver with an analytic step length. The inner solver is run under `torch.no_grad()`, so the optimized target filter is treated as a pseudo-label rather than as a fully differentiable inner optimization process.

After computing the target filter, the memory filters predicted by the memory head are trained to match it. The final auxiliary loss is the mean squared difference between each predicted filter and the optimized target filter:

$$L_{\text{match}} = w_{\text{match}} \frac{1}{M} \sum_i \|f_i - f^*\|_2^2.$$

where:

- f_i is the predicted memory filter from frame i .
- f^* is the optimized sequence-level target filter.
- w_{match} is `FILTER_MATCH_WEIGHT`.
- M is the number of predicted memory filters.

This loss has several advantages over the sampled discriminative filter loss described in Section 5.2. It provides the memory branch with a more stable target, because f^* is estimated from the full sequence rather than a small random sample. It trades the computation of cross-correlation on every frame during loss evaluation for the upfront optimization of a single sequence-level correlation filter. While running the inner optimization solver is a heavier operation, it ultimately becomes faster than the quadratic complexity emerging from recomputing cross-correlations between every predicted filter and all previous frames. Furthermore, the loss encourages temporal consistency, as all memory filters are pulled toward the same sequence-level target model.

Despite its computational advantages, the target-match loss introduces notable trade-offs. Primarily, it relies heavily on the quality of the pseudo-label. If the inner optimization yields an inaccurate filter f^* —such as during heavy occlusion—the memory tokens are penalized for imitating a flawed model. Furthermore, this supervision is highly prescriptive. By forcing memory tokens to project into a spatial correlation filter, it restricts the transformer from freely encoding other potentially useful temporal cues, such as motion or abstract relations.

Because of the architectural complexity and constrained nature of explicit auxiliary losses, the project also explored a complementary, implicit regularization strategy: Template Blurring, detailed in the following section.

5.4 Template Blurring

To complement the explicit auxiliary losses, the thesis also explores **template blurring** as an implicit memory-supervision strategy. Because the baseline OTrack architecture already possesses a highly effective template-search matching pathway, the model can often solve the training task without utilizing the newly introduced memory tokens. Template blurring artificially degrades the reliability of the initial template, forcing the tracker to rely on the recurrent memory tokens as an alternative source of target appearance information.

Let z denote the original template and $\tilde{z}$ the corrupted template. Two corruption modes were considered: a blur mode ($\tilde{z} = \text{Blur}(z)$) using repeated average pooling, and a zero-masking mode ($\tilde{z} = 0$). The blur mode is less aggressive, removing high-frequency appearance details while preserving coarse structural cues. However, because the model might still recover useful information from a blurred template, zero-masking was also introduced. By completely removing the template’s visual features, zero-masking creates maximum pressure on the network to extract and utilize the historical target information stored in the memory tokens.

During training, template corruption is applied randomly to a subset of frames within the sampled sequence. A critical design choice is that the first few search frames are strictly excluded from corruption. This allows the tracker to process clean frames ini-

tially and build a reliable memory state before template corruption is introduced. Consequently, this mechanism is only applied to sufficiently long training sequences to ensure a proper balance between clean and corrupted inputs.

Template corruption is strictly a training-time regularization technique; during inference, the tracker always receives the clean initial template.

The primary advantage of template blurring is its architectural simplicity. It requires no additional loss heads, inner optimization solvers, or pseudo-labels. The tracker is optimized entirely through the standard localization losses, while the input corruption naturally shifts the network’s internal reliance toward the memory stream. Furthermore, because it alters the actual tracking pathway rather than just an intermediate representation, improved performance on corrupted frames suggests that the memory tokens may provide target appearance cues.

As a result, template blurring serves as a controlled training condition that weakens the dominant template matching, offering a simple and implicit alternative to explicit auxiliary losses.

6 Experimental Setup

6.1 Dataset and protocol

The experiments are performed on the GOT-10k dataset. GOT-10k is a generic object tracking benchmark with more than 10,000 video segments and more than 1.5 million labeled bounding boxes [20]. The test annotations are hidden, and evaluation is performed through the official benchmark server. The tracking training in this thesis uses GOT-10k only, without additional tracking datasets.

6.2 Model configuration

As a baseline, the OTrack model with a 256-pixel search crop and Candidate Elimination (CE) pruning was chosen. Following the original OTrack configuration, the ViT backbone is initialized from an MAE-pretrained checkpoint.

The proposed MemOTrack configuration builds upon this baseline with the following specifications. The model processes a 128×128 pixel template crop (64 template tokens) and a 256×256 pixel search crop (256 search tokens). To this visual sequence, 64 memory tokens are added, bringing the initial sequence length to 384 tokens. Memory tokens are updated inside the transformer backbone via the proposed two-stage GRU recurrent mechanism.

The training pipeline employs a causal consecutive sequence sampler under a two-stage curriculum. To stabilize early optimization, the first 40 epochs utilize ground-truth-centered crops with random jitter over 20-frame rollouts. After this stage, inference-like dynamic cropping is fully enabled, and the rollout length is reduced to 15 frames.

6.3 Hyperparameter and optimizer considerations

The memory-augmented model has two different kinds of trainable parameters. The pretrained backbone already contains general visual representations. The memory tokens, GRU modules, and memory embeddings are new or substantially changed. It is therefore reasonable to use different optimizer parameter groups. The backbone can be trained with a conservative learning rate, while memory-specific parameters can receive a separately controlled learning rate.

This distinction was important in the project because the newly introduced memory tokens, recurrent update modules, and memory-read embeddings had to be detected and grouped correctly. If they were accidentally treated as ordinary pretrained backbone parameters, they might receive a learning rate too small for new modules. If they were treated too aggressively, they could destabilize the backbone. The correct choice is an empirical hyperparameter, but the optimizer must at least expose the distinction.

Gradient clipping is another relevant hyperparameter. Recurrent sequence training accumulates losses over multiple frames and backpropagates through a memory chain. Dynamic cropping can also create hard examples when early predictions move the crop away from the target. These factors can lead to occasional large gradients. Gradient clipping is therefore a practical stabilization technique, especially when newly initialized recurrent parameters are trained together with a pretrained transformer backbone [28].

6.4 Evaluation protocol

The tracker is evaluated on the GOT-10k test set using Average Overlap (AO), Success Rate at 0.50 (SR0.50), and Success Rate at 0.75 (SR0.75), as defined in Section 1.4.

The comparison is made against reported OTrack-256 + CE and OTrack-384 + CE baselines. The 256 baseline is the closest architectural comparison because it uses the same search resolution. The 384 baseline is included as a stronger high-resolution reference. Because the proposed MemOTrack model introduces several changes simultaneously, the final result should not be interpreted as identifying a single cause. A complete study would include ablations such as memory tokens without dynamic cropping, dynamic cropping without memory, one GRU instead of two GRUs, and different memory

loss weights. These experiments are left for future work.

7 Results and Discussion

7.1 Quantitative results

The final measured result for MemOTrack-256 + CE is shown in Table 7.1. The model achieved AO 0.729, SR0.50 0.823, and SR0.75 0.693. The reported OTrack-256 + CE baseline achieves AO 0.710, SR0.50 0.804, and SR0.75 0.682. The reported OTrack-384 + CE baseline achieves AO 0.737, SR0.50 0.832, and SR0.75 0.708 [13].

Table 7.1 - Final quantitative comparison on GOT-10k.

Method	AO	SR0.50	SR0.75
OTrack-256 + CE	0.710	0.804	0.682
MemOTrack-256 + CE	0.729	0.823	0.693
OTrack-384 + CE	0.737	0.832	0.708

The final configuration improves upon the reported same-resolution OTrack-256 + CE baseline across all three metrics, with gains of +0.019 in AO, +0.019 in SR0.50, and +0.011 in SR0.75. Because only one final configuration was fully evaluated, this result should be interpreted as preliminary evidence for the effectiveness of the complete MemOTrack training pipeline, rather than isolated proof that the recurrent memory alone caused the gain.

However, the results remain below those of OTrack-384 + CE, with gaps of -0.008 in AO, -0.009 in SR0.50, and -0.015 in SR0.75. This indicates that while the memory-augmented 256-resolution model closes the gap to the stronger high-resolution baseline, it does not surpass it.

7.2 Training Dynamics

To understand how the model adapts over time, the learning dynamics were analyzed throughout the training process. Figure 7.1 and Figure 7.2 illustrate the training and validation metrics, respectively, which clearly reflect the distinct phases of the training strategy.




Figure 7.1 - Training loss and overlap metrics over the 75 recorded epochs, illustrating the impact of the different training phases.

The training curve exhibits rapid early convergence during the initial warm-up phase (epochs 1–20) under standard ground-truth-centered cropping. In this period, the total training loss decreases sharply from 2.033 to 0.547, while the training IoU climbs to approximately 0.870.

Between epochs 21 and 40, template blurring was introduced to the pipeline. During this phase, the standard training IoU remained stable and high. More importantly, the model’s performance on specifically corrupted templates—tracked via the separate Blur/IoU metric—steadily improved from 0.849 at epoch 21 to a peak of 0.875 at epoch 39. This upward trend indicates that the network adapted to the degraded initial tem-

plate, which may have encouraged greater use of the recurrent memory state to maintain accurate localization.

At epoch 41, template blurring was disabled, and the training pipeline transitioned to the inference-like dynamic rollout setting. This shift intentionally introduced a significantly harder optimization objective, as the model was now forced to handle its own prediction drift across frames. This increased difficulty is immediately visible in the graphs: the GIoU loss spikes from 0.114 to 0.139, and the training IoU drops from 0.891 to 0.869. However, the model demonstrated robust recovery in the subsequent epochs, successfully adapting to the dynamic crops and restoring its IoU to previous levels by epoch 59.

A scheduled learning rate decay at epoch 60 resulted in a further, immediate reduction in training loss (dropping from 0.390 to 0.364) and a corresponding jump in training accuracy.




Figure 7.2 - Validation loss and overlap metrics recorded every five epochs.

The validation metrics (Figure 7.2) corroborate the trends observed during training, confirming that the model did not overfit when transitioning to the harder dynamic rollout phase. The validation loss reached its lowest point of 0.442 at epoch 60. Although

the internal validation IoU continued to rise slightly in later epochs (reaching 0.885 at epoch 75), the epoch-60 checkpoint exhibited the best overall validation loss profile. Consequently, this checkpoint was selected for the official, external GOT-10k test evaluation reported in Table 7.1. This distinction is important: the local validation curves describe internal training dynamics, whereas the final benchmark results are derived from the optimal checkpoint identified by these curves.

7.3 Computational Cost and Token Budget

The performance differences between the models can be further contextualized by analyzing the token budget processed by the ViT backbone, which serves as a proxy for computational complexity.

As shown in Table 7.2, the standard OTrack-256 setting produces an initial sequence of 320 visual tokens. MemOTrack-256 + CE adds 64 learned memory tokens, increasing the initial sequence to 384 tokens. In contrast, the OTrack-384 baseline processes a much larger 192×192 template and 384×384 search region, resulting in 720 initial visual tokens.

Table 7.2 - Input token counts before candidate elimination.

Model setting	Template tokens	Search tokens	Memory tokens	Initial tokens
OTrack-256 + CE	64	256	0	320
MemOTrack-256 + CE	64	256	64	384
OTrack-384 + CE	144	576	0	720
MemOTrack-384 + CE, derived	144	576	64	784

Candidate elimination (CE) dynamically reduces the active sequence length inside the backbone. Table 7.3 illustrates the token counts after the three CE pruning stages (using a keep ratio of 0.7). Because memory tokens are preserved across CE layers,

MemOSTrack-256 retains 217 active tokens in the final layers, compared to 153 for the baseline OTrack-256.

Table 7.3 - Active token counts after CE pruning.

Model setting	Initial tokens	After CE 1	After CE 2	After CE 3
OTrack-256 + CE	320	244	190	153
MemOSTrack-256 + CE	384	308	254	217
OTrack-384 + CE	720	548	427	343

This token-budget analysis clarifies the quantitative results. MemOSTrack-256 + CE operates with a 20% larger initial sequence than its direct baseline, yielding measurable accuracy improvements. However, it remains computationally much lighter than OTrack-384 + CE, which processes nearly double the number of tokens in the final layers (343 vs. 217). The remaining accuracy gap to the 384-resolution model is therefore consistent with the significant difference in their spatial token budgets.

7.4 Limitations

The most important limitation is the lack of a full ablation study. Because multiple changes were introduced, it is not possible to isolate the exact effect of each component from the final table alone. For example, the gain over OTrack-256 + CE may come from memory tokens, dynamic cropping, candidate elimination interactions, training differences, template blurring, auxiliary losses, or a combination of these factors.

The second limitation is that the memory mechanism was evaluated mainly on the standard metrics. It may be more useful to analyze specific sequence attributes such as occlusion, deformation, or long-term appearance change. If memory helps only in certain cases, the average score may hide this benefit.

The third limitation is that memory loss heads were not fully optimized. The balance

between the main loss and memory loss is difficult and likely important. A scheduled memory loss or a warm-up phase may be more stable than applying the full auxiliary loss from the beginning.

7.5 Future work

Future work should begin with ablation experiments. The first ablation should compare baseline OTrack, OTrack with sequence training but no memory, MemOTrack with static ground-truth-centered crops, and MemOTrack with dynamic crops. This would separate the effects of memory and cropping.

The second ablation should compare memory update mechanisms. A single temporal GRU, the current two-stage GRU, and the three-way softmax gated update should be tested under the same training conditions. This would reveal whether the layer-wise recurrent connection is helpful.

The third ablation should vary the memory supervision strategy. Possible variants include no memory head, memory head after selected layers only, scheduled memory loss, and different values of `lambda_mem`. Template blurring should also be varied by blur strength and probability.

Finally, the model should be evaluated on sequences where memory is expected to matter: long occlusions, viewpoint changes, and gradual appearance changes. If memory helps only on a subset of videos, attribute-level analysis may reveal benefits that average metrics hide.

Conclusion

This thesis investigated the integration of recurrent memory tokens into OTrack. The proposed MemOTrack model extends the original one-stream transformer tracker with memory tokens and a two-stage GRU update mechanism. The first GRU combines current candidate memory with memory from the previous frame, while the second GRU combines the result with memory from the previous transformer layer. This design was intended to provide both temporal continuity and layer-wise memory refinement.

The work also required a substantial rewrite of the training pipeline. The original pair-based sampling strategy was replaced by sequence-based sampling so that memory could be propagated through ordered video frames. The final version further introduced dynamic search cropping, where the search crop is generated from the previous prediction rather than from the current ground truth. This made training more similar to inference, although also more difficult.

Auxiliary memory loss heads were explored because a strong OTrack baseline can ignore newly added memory tokens. Template blurring was considered as a complementary solution to reduce over-reliance on the clean initial template and encourage use of recurrent memory.

The final MemOTrack-256 + CE model achieved AO 0.729, SR0.50 0.823, and SR0.75 0.693. This improves on the reported same-resolution OTrack-256 + CE baseline, but it does not outperform the stronger higher-resolution OTrack-384 + CE baseline. The result provides a useful foundation for future work on memory supervision, update mechanisms, resolution scaling, and training curricula for transformer-based visual tracking.

References

- [1] D. Comaniciu, V. Ramesh, and P. Meer, ‘Real-Time Tracking of Non-Rigid Objects Using Mean Shift,’ CVPR, 2000.
- [2] B. D. Lucas and T. Kanade, ‘An Iterative Image Registration Technique with an Application to Stereo Vision,’ IJCAI, 1981.
- [3] M. Isard and A. Blake, ‘CONDENSATION - Conditional Density Propagation for Visual Tracking,’ International Journal of Computer Vision, 1998.
- [4] I. Matthews, T. Ishikawa, and S. Baker, ‘The Template Update Problem,’ IEEE Transactions on Pattern Analysis and Machine Intelligence, 2004.
- [5] D. S. Bolme, J. R. Beveridge, B. A. Draper, and Y. M. Lui, ‘Visual Object Tracking using Adaptive Correlation Filters,’ CVPR, 2010.
- [6] J. F. Henriques, R. Caseiro, P. Martins, and J. Batista, ‘Exploiting the Circulant Structure of Tracking-by-Detection with Kernels,’ ECCV, 2012.
- [7] J. F. Henriques, R. Caseiro, P. Martins, and J. Batista, ‘High-Speed Tracking with Kernelized Correlation Filters,’ IEEE Transactions on Pattern Analysis and Machine Intelligence, 2015.
- [8] M. Danelljan, G. Hager, F. Shahbaz Khan, and M. Felsberg, ‘Accurate Scale Estimation for Robust Visual Tracking,’ BMVC, 2014.
- [9] M. Danelljan, G. Bhat, F. Shahbaz Khan, and M. Felsberg, ‘ECO: Efficient Convolution Operators for Tracking,’ CVPR, 2017.
- [10] L. Bertinetto, J. Valmadre, J. F. Henriques, A. Vedaldi, and P. H. S. Torr, ‘Fully-

Convolutional Siamese Networks for Object Tracking,’ ECCV Workshops, 2016.

[11] B. Li, J. Yan, W. Wu, Z. Zhu, and X. Hu, ‘High Performance Visual Tracking with Siamese Region Proposal Network,’ CVPR, 2018.

[12] B. Yan, H. Peng, J. Fu, D. Wang, and H. Lu, ‘Learning Spatio-Temporal Transformer for Visual Tracking,’ ICCV, 2021.

[13] B. Ye, H. Chang, B. Ma, S. Shan, and X. Chen, ‘Joint Feature Learning and Relation Modeling for Tracking: A One-Stream Framework,’ ECCV, 2022.

[14] V. Borsuk, R. Vei, O. Kupyn, T. Martyniuk, I. Krasheniyi, and J. Matas, ‘FEAR: Fast, Efficient, Accurate and Robust Visual Tracker,’ ECCV, 2022.

[15] Y. Zheng, B. Zhong, Q. Liang, Z. Mo, S. Zhang, and X. Li, ‘ODTrack: Online Dense Temporal Token Learning for Visual Tracking,’ AAAI, 2024.

[16] Y. Bai, Z. Zhao, Y. Gong, and X. Wei, ‘ARTrackV2: Prompting Autoregressive Tracker Where to Look and How to Describe,’ CVPR, 2024.

[17] S. Liang, Y. Bai, Y. Gong, and X. Wei, ‘Autoregressive Sequential Pretraining for Visual Tracking,’ CVPR, 2025.

[18] G. Bhat, M. Danelljan, L. Van Gool, and R. Timofte, ‘Learning Discriminative Model Prediction for Tracking,’ ICCV, 2019.

[19] M. Danelljan, L. Van Gool, and R. Timofte, ‘Probabilistic Regression for Visual Tracking,’ CVPR, 2020.

[20] L. Huang, X. Zhao, and K. Huang, ‘GOT-10k: A Large High-Diversity Benchmark for Generic Object Tracking in the Wild,’ IEEE Transactions on Pattern Analysis and Machine Intelligence, 2021.

[21] A. Vaswani et al., ‘Attention Is All You Need,’ NeurIPS, 2017.

[22] A. Dosovitskiy et al., ‘An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale,’ ICLR, 2021.

[23] K. Cho et al., ‘Learning Phrase Representations using RNN Encoder-Decoder for

Statistical Machine Translation,’ EMNLP, 2014.

[24] K. He, X. Chen, S. Xie, Y. Li, P. Dollar, and R. Girshick, ‘Masked Autoencoders Are Scalable Vision Learners,’ CVPR, 2022.

[25] M. Kim, S. Lee, J. Ok, B. Han, and M. Cho, ‘Towards Sequence-Level Training for Visual Tracking,’ ECCV, 2022.

[26] A. Zhang, Z. C. Lipton, M. Li, and A. J. Smola, ‘Vision Transformer.svg,’ Wikimedia Commons, 2024. Licensed under Creative Commons Attribution-ShareAlike 4.0 International (CC BY-SA 4.0). Available: https://commons.wikimedia.org/wiki/File:Vision_Transformer.svg

[27] R. J. Williams and D. Zipser, ‘Gradient-Based Learning Algorithms for Recurrent Networks and Their Computational Complexity,’ in Backpropagation: Theory, Architectures, and Applications, 1995.

[28] R. Pascanu, T. Mikolov, and Y. Bengio, ‘On the Difficulty of Training Recurrent Neural Networks,’ Proceedings of the 30th International Conference on Machine Learning (ICML), 2013.

Abstract

Extending OTrack with GRU-Based Memory Tokens for Visual Object Tracking

This thesis investigates the integration of recurrent memory tokens into OTrack, a one-stream transformer tracker. The proposed MemOTrack model adds memory tokens to the template-search token sequence and updates them with a two-stage GRU mechanism using previous-frame and previous-layer memory. The training pipeline was rewritten from pair-based sampling to ordered video sequence sampling and further extended with dynamic search cropping to better match inference. Auxiliary memory loss heads and template blurring were explored to encourage the model to use memory. The final model achieved AO 0.729, SR0.50 0.823, and SR0.75 0.693. It exceeded the reported same-resolution OTrack-256 + CE baseline, but remained slightly below the stronger reported OTrack-384 + CE baseline.

Keywords: visual object tracking, OTrack, transformer, GRU, memory tokens, dynamic cropping

Summary

Extending OTrack with GRU-Based Memory Tokens for Visual Object Tracking

The thesis studies whether a transformer-based tracker can benefit from explicit recurrent memory. OTrack is selected as the baseline because it is a strong one-stream tracker that jointly processes template and search tokens. The proposed MemOTrack model inserts memory tokens into this token sequence. A transformer layer first produces candidate memory, after which a two-stage GRU update combines the candidate with previous-frame memory and previous-layer memory.

A major part of the work is the training pipeline. The original pair-based sampler was not sufficient because recurrent memory needs ordered video frames. Therefore, the sampler was rewritten to return causal consecutive search sequences. The final pipeline also uses dynamic search cropping, where each next search crop is produced from the previous prediction. This makes training closer to real inference.

The thesis also discusses memory loss heads and template blurring. Memory loss heads were introduced to prevent the model from ignoring memory tokens. Template blurring was explored to reduce over-reliance on the clean initial template and encourage the use of temporal memory.

The final evaluation shows that MemOTrack-256 + CE achieved AO 0.729, SR0.50 0.823, and SR0.75 0.693. These results are higher than the reported OTrack-256 + CE baseline, but lower than the reported OTrack-384 + CE baseline. The main conclusion is that the final MemOTrack configuration improves the same-resolution 256 CE setting, but does not yet surpass the higher-resolution 384 CE model. Further optimization

and ablation studies are required to separate the effects of memory, dynamic cropping, candidate elimination, and resolution.

Abbreviations

Table A.1 - Abbreviations.

Abbreviation English term		Explanation
AO	Average Overlap	Average IoU over evaluated frames
BPTT	Backpropagation Through Time	Training procedure for recurrent models unrolled over sequences
CE	Candidate Elimination	OTrack efficiency module that removes unlikely tokens
GOT-10k	Generic Object Tracking benchmark	Large-scale tracking benchmark
GRU	Gated Recurrent Unit	Recurrent neural network unit with gates
IoU	Intersection over Union	Overlap between predicted and ground-truth boxes
MAE	Masked Autoencoder	Self-supervised vision pretraining method
SOT	Single-Object Tracking	Tracking one specified object through a video
SR	Success Rate	Fraction of frames above an IoU threshold
ViT	Vision Transformer	Transformer model applied to image patches